

## Sequence to Sequence modelling(HMM)

Hidden Markov Models (HMMs) are incredibly versatile tools for modeling sequential data where we deal with hidden states that generate observable events.

---

### Key Concepts of HMM

1. **Hidden States:**
    - Represent the unobservable states of the system (e.g., "Sunny" or "Rainy").
    - Transition probabilities govern the likelihood of moving between these states.
  2. **Observations:**
    - Represent observable outcomes generated by the hidden states (e.g., "Dry" or "Wet").
    - Emission probabilities define the likelihood of observing a specific event given the current hidden state.
  3. **Core Probabilities:**
    - **Initial State Distribution:** Probability of starting in each hidden state.
    - **Transition Probabilities:** Probability of moving from one hidden state to another.
    - **Emission Probabilities:** Probability of observing a specific event given a hidden state.
- 

### HMM Algorithm Workflow

1. **Define Model Parameters:**
    - State space, observation space, initial probabilities, transition matrix, and emission matrix.
  2. **Training the Model:**
    - Use algorithms like Baum-Welch (or Forward-Backward) to estimate model parameters from data.
  3. **Decoding:**
    - Use the Viterbi algorithm to find the most likely sequence of hidden states given observations.
  4. **Evaluation:**
    - Validate the model using metrics such as accuracy or likelihood scores.
- 

### Python Implementation Examples

#### Example 1: Weather Prediction

Code:

#### Step 1: Import Required Libraries

```
import numpy as np
from hmmlearn import hmm
import matplotlib.pyplot as plt
import seaborn as sns
```

### Explanation:

- **numpy:** Used for numerical computations like creating probability matrices and observation sequences.
  - **hmmlearn:** Provides tools to implement Hidden Markov Models (HMMs). In this example, we use `hmm.CategoricalHMM` to define and work with discrete observation spaces.
  - **matplotlib and seaborn:** Used to visualize results like predicted hidden states.
- 

## Step 2: Define States and Observations

```
# Define states and observations
states = ["Sunny", "Rainy"]
observations = ["Dry", "Wet"]
n_states = len(states)
n_observations = len(observations)
```

### Explanation:

- **states:** The hidden states of the model (Sunny and Rainy), which are not directly observable.
  - **observations:** The measurable events (Dry and Wet) that depend on the hidden states.
  - **n\_states and n\_observations:** Used to specify the number of states and observations in the model.
- 

## Step 3: Define Model Parameters

```
# Define model parameters
state_probability = np.array([0.6, 0.4]) # Initial probabilities
transition_probability = np.array([[0.7, 0.3], [0.3, 0.7]]) # Transition
matrix
emission_probability = np.array([[0.9, 0.1], [0.2, 0.8]]) # Emission
matrix
```

### Explanation:

- **state\_probability:**
  - Initial probabilities for each state:
    - 60% chance of starting in Sunny.
    - 40% chance of starting in Rainy.
- **transition\_probability:**
  - A **2x2 matrix** describing the likelihood of transitioning between states:

- If it's `Sunny`, there's a 70% chance it remains `Sunny` and a 30% chance it becomes `Rainy`.
  - If it's `Rainy`, there's a 70% chance it remains `Rainy` and a 30% chance it becomes `Sunny`.
  - **emission\_probability:**
    - A **2x2 matrix** describing the likelihood of observations given the states:
      - If it's `Sunny`, there's a 90% chance of observing `Dry` and a 10% chance of observing `Wet`.
      - If it's `Rainy`, there's a 20% chance of observing `Dry` and an 80% chance of observing `Wet`.
- 

## Step 4: Initialize the HMM Model

```
# Create HMM model
model = hmm.CategoricalHMM(n_components=n_states)
model.startprob_ = state_probability
model.transmat_ = transition_probability
model.emissionprob_ = emission_probability
```

### Explanation:

- **hmm.CategoricalHMM(n\_components=n\_states):**
    - Creates an HMM with discrete observation spaces (categories).
    - `n_components` specifies the number of hidden states (`Sunny`, `Rainy`).
  - **model.startprob\_:**
    - Sets the initial state probabilities.
  - **model.transmat\_:**
    - Sets the state transition probabilities.
  - **model.emissionprob\_:**
    - Sets the emission probabilities.
- 

## Step 5: Define Observation Sequence

```
# Define sequence of observations
observations_sequence = np.array([0, 1, 0, 1, 0, 0]).reshape(-1, 1)
```

### Explanation:

- **observations\_sequence:**
  - Represents observed events over time:
    - 0: "Dry"
    - 1: "Wet"
  - The sequence `[0, 1, 0, 1, 0, 0]` means:
    - Time step 1: Dry
    - Time step 2: Wet
    - Time step 3: Dry, and so on.
- **.reshape(-1, 1):**

- Reshapes the sequence into a column vector, as required by `hmmlearn`.
- 

## Step 6: Predict the Most Likely Sequence of Hidden States

```
# Predict hidden states
hidden_states = model.predict(observations_sequence)
```

### Explanation:

- `model.predict()`:
    - Uses the HMM model to predict the **hidden state sequence** for the given observations.
    - Returns the most likely sequence of hidden states based on the input observations.
- 

## Step 7: Decode Observations with Viterbi Algorithm

```
# Decode observations using Viterbi algorithm
log_prob, viterbi_hidden_states = model.decode(observations_sequence,
algorithm="viterbi")
```

### Explanation:

- `model.decode()`:
    - Uses the Viterbi algorithm to compute:
      - **Log Probability (`log_prob`)**: Likelihood of the decoded hidden state sequence.
      - **Viterbi Hidden States (`viterbi_hidden_states`)**: The most probable sequence of hidden states.
- 

## Step 8: Display Results

```
print("Most likely hidden states:", hidden_states)
print("Viterbi Log Probability:", log_prob)
```

### Explanation:

- Displays the predicted hidden states (`hidden_states`) and the log probability of the sequence (`log_prob`).
- 

## Step 9: Visualize Results

```
# Plotting
```

```
sns.set_style("whitegrid")
plt.plot(hidden_states, '-o', label="Predicted Hidden State")
plt.xlabel("Time Step")
plt.ylabel("Hidden State (Sunny=0, Rainy=1)")
plt.title("Weather Prediction (Sunny or Rainy)")
plt.legend()
plt.show()
```

### Explanation:

- **Plot:**
  - X-axis: Time steps.
  - Y-axis: Predicted hidden states (0 = Sunny, 1 = Rainy).
  - Visualization provides an intuitive view of how the model predicts state transitions over time.

The **Viterbi algorithm** is a dynamic programming algorithm used to find the most probable sequence of hidden states (also called the **optimal path**) in a **Hidden Markov Model (HMM)**, given a sequence of observations. It is a core method in applications like speech recognition, natural language processing, and bioinformatics.

When working with HMMs, we often want to infer the **hidden state sequence** that most likely explains the observed sequence. The Viterbi algorithm efficiently computes this sequence by:

1. Maximizing the probability of the entire path, rather than individual states.
2. Avoiding a brute-force computation of all possible paths (which would be computationally expensive).

---

### Example 2: Speech Recognition step-by-step.

---

#### Step 1: Import Required Libraries

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from hmmlearn import hmm
```

### Explanation:

- The libraries here are identical to the weather prediction example:
    - **numpy**: Numerical computations.
    - **hmmlearn**: Provides the `CategoricalHMM` class to work with discrete observations.
    - **matplotlib and seaborn**: For visualizing the predicted hidden states.
-

## Step 2: Define States and Observations

```
# Define the state space
states = ["Silence", "Word1", "Word2", "Word3"]
n_states = len(states)

# Define the observation space
observations = ["Loud", "Soft"]
n_observations = len(observations)
```

### Explanation:

- **states:**
    - Hidden states in this model represent:
      - "Silence" (no speech).
      - Three distinct words: Word1, Word2, and Word3.
    - These states are not directly observable.
  - **observations:**
    - The measurable events ("Loud" and "Soft") are volume levels of the speech.
- 

## Step 3: Define Model Parameters

```
# Define the initial state distribution
start_probability = np.array([0.8, 0.1, 0.1, 0.0])

# Define the state transition probabilities
transition_probability = np.array([[0.7, 0.2, 0.1, 0.0],
                                   [0.0, 0.6, 0.4, 0.0],
                                   [0.0, 0.0, 0.6, 0.4],
                                   [0.0, 0.0, 0.0, 1.0]])

# Define the observation likelihoods
emission_probability = np.array([[0.7, 0.3],
                                   [0.4, 0.6],
                                   [0.6, 0.4],
                                   [0.3, 0.7]])
```

### Explanation:

- **Initial State Probabilities (start\_probability):**
  - High likelihood (80%) of starting in the "Silence" state.
  - Lower probabilities (10%) of starting in Word1 or Word2. No chance of starting in Word3.
- **Transition Probabilities (transition\_probability):**
  - Defines the likelihood of moving between states:
    - "Silence" tends to remain silent (70%) but can transition to Word1 (20%) or Word2 (10%).
    - Word1 can transition to Word2 (40%) or remain in Word1 (60%).
    - Word2 can transition to Word3 (40%) or remain in Word2 (60%).
    - Word3 always remains in Word3 (100%).
- **Emission Probabilities (emission\_probability):**
  - Defines the likelihood of observing "Loud" or "Soft" given a hidden state:

- "Silence" is more likely to produce "Loud" (70%).
  - Word1 and Word2 have mixed likelihoods of emitting "Loud" or "Soft".
  - Word3 is more likely to produce "Soft" (70%).
- 

## Step 4: Initialize the HMM Model

```
# Fit the model
model = hmm.CategoricalHMM(n_components=n_states)
model.startprob_ = start_probability
model.transmat_ = transition_probability
model.emissionprob_ = emission_probability
```

### Explanation:

- The HMM model is created using the `hmm.CategoricalHMM` class with:
    - **n\_components=n\_states**: Specifies 4 hidden states (Silence, Word1, Word2, Word3).
  - The model parameters are set:
    - **Initial probabilities**: `start_probability`.
    - **Transition probabilities**: `transition_probability`.
    - **Emission probabilities**: `emission_probability`.
- 

## Step 5: Define Observation Sequence

```
# Define the sequence of observations
observations_sequence = np.array([0, 1, 0, 0, 1, 1, 0, 1]).reshape(-1, 1)
```

### Explanation:

- **observations\_sequence**:
    - Represents observed events over time:
      - 0: "Loud"
      - 1: "Soft"
    - The sequence `[0, 1, 0, 0, 1, 1, 0, 1]` corresponds to:
      - Time step 1: Loud
      - Time step 2: Soft
      - Time step 3: Loud, and so on.
- 

## Step 6: Predict the Most Likely Sequence of Hidden States

```
# Predict the most likely hidden states
hidden_states = model.predict(observations_sequence)
print("Most likely hidden states:", hidden_states)
```

### Explanation:

- `model.predict()`:
    - Computes the **most probable sequence of hidden states** that explains the given observations.
  - **Output (`hidden_states`):**
    - A sequence of integers corresponding to the hidden states:
      - 0: "Silence"
      - 1: "Word1"
      - 2: "Word2"
      - 3: "Word3"
- 

## Step 7: Visualize Results

```
# Plot the results
sns.set_style("darkgrid")
plt.plot(hidden_states, '-o', label="Predicted Hidden State")
plt.xlabel("Time Step")
plt.ylabel("Hidden State")
plt.title("Speech Recognition Hidden States")
plt.legend()
plt.show()
```

### Explanation:

- The plot shows how the model predicts hidden states over time.
  - **X-axis:** Time steps.
  - **Y-axis:** Hidden states (Silence, Word1, Word2, Word3).
- 

## Output Example

### Predicted Hidden States:

```
Most likely hidden states: [0 1 2 2 3 3 3 3]
```

### Interpretation:

- **Time Step 1:** "Silence".
- **Time Step 2:** Transitions to "Word1".
- **Time Steps 3 & 4:** "Word2".
- **Time Steps 5-8:** "Word3".

### Visualization:

- The graph reflects these transitions, illustrating the most likely hidden states over time.
- 

## Applications in Speech Recognition



1. **State Space:** Represents phonemes, words, or speech segments (e.g., Silence, Word 1, etc.).
2. **Observation Space:** Represents features of speech (e.g., volume, frequency).
3. **Use Cases:**
  - Identify words or phrases spoken in audio data.
  - Segment audio into meaningful parts (e.g., silence vs. spoken words).